

Vedant Vardhaan

San Diego, CA | 858-752-3541 | vvardhaan@ucsd.edu

[linkedin.com/in/vedantvardhaan](https://www.linkedin.com/in/vedantvardhaan) | github.com/VedVar43789 | vedantvardhaan23.vercel.app

EDUCATION

University of California, San Diego

San Diego, CA

Bachelor of Science in Data Science, Minor in Finance | **GPA: 3.96**

Sep 2023 – Jun 2027

- **Relevant Coursework:** Data Structures, Algorithms, Machine Learning, Deep Learning, Object-Oriented Programming

EXPERIENCE

UniCircle

San Diego, CA

Chief Technology Officer and Full-Stack Engineer

May 2026 – Present

- Own unicircle.app end to end as founding engineer at a UCSD incubator-listed startup, shipping React and TypeScript product surfaces from database schema through production deployment for a club discovery platform.
- Designed typed REST interfaces between the front end and a PostgreSQL and Supabase Edge Functions backend, owning end-to-end data workflows with row-level security enforcing FERPA and CCPA compliance on all student-data access.
- Set code review, testing, and CI/CD standards for all contributors and advise the CEO on technical roadmap. **Tech:** React, TypeScript, Supabase, PostgreSQL, Vercel.

Pricewaterhouse Coopers LLP (PwC)

Gurugram, India

Technology Consultant Intern, AI Engineering

Jun 2025 – Sep 2025

- Cut cloud provisioning time by 60% by integrating modern LLMs into a production internal developer platform, a LangGraph and Gemini 2.5 Pro multi-agent system orchestrating AWS and Azure onboarding across 15+ services.
- Shipped modular Python backend services and REST endpoints for credential management, IAM configuration, and resource provisioning, invoked directly by engineering teams.
- Reduced manual configuration errors by 45% by instrumenting infrastructure pipelines with Azure Key Vault, CloudTrail, and GuardDuty for auditing and observability. **Tech:** LangGraph, LangChain, Python, AWS, Azure.

CSES Open Source, UC San Diego

San Diego, CA

President and Engineering Manager

Nov 2024 – Present

- Direct a 100-member open-source organization, launching 3 projects and partnering with the Linux Foundation, Cisco, and NVIDIA.
- As Engineering Manager of TritonSpend, ship a mobile-first financial wellness app adopted by 100+ UCSD students on real-time transaction and budgeting APIs.
- Lead a 6-engineer team through agile sprints, design review, and CI/CD, cutting deployment bugs by 40% and improving delivery speed by 30%. **Tech:** React Native, TypeScript, Node.js, PostgreSQL, Supabase.

Engineers for Exploration, UC San Diego

San Diego, CA

Undergraduate Researcher and Project Lead

Oct 2025 – Apr 2026

- Led a geospatial data ingestion pipeline fusing drone and Sentinel-2 imagery across 10+ GB of remote-sensing data, improving model accuracy by 25% over baseline, and delivered production inference workflows to Scripps Institution of Oceanography. **Tech:** Python, TensorFlow, Google Earth Engine.

PROJECTS

TumorSense: Clinical Imaging Diagnosis | *PyTorch, Scikit-learn, FastAPI, RAG* | [GitHub](#) Apr 2026 – Jun 2026

- Built a dual-track medical imaging classifier combining SVM with a CNN and Vision Transformer ensemble, reaching 98% malignant recall and 0.995 ROC-AUC under a rigorous evaluation harness; awarded **Best Project, DS3 Spring 2026**.
- Deployed model inference through a FastAPI service with SHAP and RAG-based explainability, surfacing per-prediction clinical rationale to support reviewer trust in high-stakes diagnostic decisions.

InvasiveWatch: Global Species Tracker | *React Native, FastAPI, FAISS, Python* | [GitHub](#) Jan 2026 – Mar 2026

- Engineered a retrieval and risk-ranking pipeline on FAISS vector search over 96K+ species, reaching 83.4% recall under an evaluation harness scoring climate, trait, and invasion-history signals; awarded **Best Project, DS3 Winter 2026**.
- Built a React Native and FastAPI platform serving model inference endpoints that surface invasive-species risk across user-selected regions from multiple biodiversity APIs.

Optimizing Qwen3-4B for Math Reasoning | *PyTorch, LLM Inference* | [GitHub](#) May 2026 – Jun 2026

- Raised Qwen3-4B math-reasoning accuracy to 0.720 through inference-time optimization of the prompt builder, fine-tuning and answer post-processing pipeline, ranking 19th (out of 120) on the leaderboard.

TECHNICAL SKILLS

Languages: Python, TypeScript, JavaScript, Java, SQL, R | **Frontend:** React, React Native, D3.js, HTML/CSS

Backend and Databases: Node.js, Express.js, FastAPI, REST APIs, PostgreSQL, Supabase, SQL

Cloud and DevOps: AWS (EC2, S3, IAM), Azure, GCP, Docker, Kubernetes, CI/CD, Git/GitHub, Vercel

AI/ML: LangChain, LangGraph, LLM Integration, LLM Evaluation, RAG, PyTorch, TensorFlow, Hugging Face,

Scikit-learn, FAISS, Pandas, NumPy | **Tools:** Claude Code, Cursor